

Master Thesis

Recommender System for Galaxy Tools and Workflows

(Find similar tools and predict next tools in workflows)

Anup Kumar

Examiners: Prof. Dr. Rolf Backofen

Prof. Dr. Wolfgang Hess

Adviser: Dr. Björn Grüning

University of Freiburg
Faculty of Engineering
Department of Computer Science
Chair for Bioinformatics

July, 2018

Thesis period

10.01.2018 – 09.07.2018

Examiners

Prof. Dr. Rolf Backofen and Prof. Dr. Wolfgang Hess

Adviser

Dr. Björn Grüning

Declaration

I hereby declare, that I am the sole author and composer of my thesis and that no other sources or learning aids, other than those listed, have been used. Furthermore, I declare that I have acknowledged the work of others by providing detailed references of said work.

I hereby also declare, that my Thesis has not been prepared for another examination or assignment, either wholly or excerpts thereof.

Place, Date

Signature

Acknowledgment

I express my gratitude to all the people who encouraged and supported me to accomplish this work. I am grateful to my mentor Dr. Björn Grüning who entrusted me with the task of building a recommendation system for Galaxy. He facilitated this work by providing me with all the indispensable means. Being specific, his pragmatic suggestions concerning the Galaxy tools and workflows helped me discern them better and improve the performance. His advice to create a visualizer for showing the similar tools worked wonders as it enabled me to find and rectify a few bugs which were tough to establish. For the next task, creating a separate visualizer for looking through the next predicted tools was conducive in all merits. I appreciate and thank Eric Rasche who extracted the workflows for me from the Galaxy Freiburg server. I offer thanks to Dr. Mehmet Tekman and Joachim Wolff for their expert feedback, insights and general advice. At length, I wish to thank all the other members of Freiburg Galaxy team for their continuous support and help.

Abstract

The study explores two concepts to devise a recommendation system for Galaxy. One idea is to find similar Galaxy tools for each tool and another is to predict a set of possible next tools in Galaxy workflows.

To find similarities among tools, we need to extract information about each tool from its attributes like name, description, input and output file types and help text. We take into account these attributes one by one and compute similarity matrices. We compute three similarity matrices, one each for input and output file, name and description and help text attributes. Each row in a similarity matrix holds similarity scores of one tool against all the other tools. These similarity scores depend on the similarity measures (jaccard index and cosine similarity) used to compute the score between a pair of tools. To combine these matrices, one simple solution is to compute an average. But, assigning equal importance weight to each matrix might be sub-optimal. To find an optimal combination, we use optimization to learn importance weights for the corresponding rows for each tool in the similarity matrices. To define a loss function for optimization, we use a true similarity value based on the similarity measures. The similarity scores are positive real numbers between 0 and 1. We take an array of 1.0 as the true value.

Next task analyzes workflows to predict a set of next tools at each stage of creating workflows. While creating workflows, it would be convenient to leaf through a set of next possible tools as a guide. It can assist the less experienced (Galaxy) users in creating workflows when they are unsure about which tools can further be connected. In addition, it can curtail the time taken in creating a workflow. To achieve that, we need to learn the connections among tools in order to be able to predict the next possible ones based on the previously connected tools. To preprocess the workflows to make them usable by downstream machine learning algorithms, we compute all the paths bridging the starting and end tools in all workflows. We follow a classification approach to predict the next tools and use LSTM (long short-term memory), a variant of recurrent neural networks. It performs well for learning long range, sequential and time-dependent data (tools connections) [1, 2]. We report the accuracy as precision.

Zusammenfassung

Contents

1	Introduction	2
1.1	Galaxy	2
1.2	Galaxy tools	3
1.3	Motivation	4
2	Approach	6
2.1	Extract data	6
2.1.1	Select tools attributes	6
2.1.2	Clean data	7
2.2	Document embeddings	11
2.2.1	Latent semantic indexing	11
2.2.2	Paragraph vectors	15
2.3	Similarity measures	17
2.3.1	Cosine similarity	17
2.3.2	Jaccard index	18
2.4	Optimization	18
3	Experiments	21
3.1	Amount of help text	21
3.2	One and two sources instead of three	21
3.3	Similarity measures	21
3.4	Latent Semantic Analysis	21
3.5	Paragraph Vector	22
3.5.1	Distributed bag-of-words	22
3.6	Gradient Descent	22
3.6.1	Learning rates	23
4	Results and Analysis	24
4.1	Latent Semantic Analysis	24
4.1.1	Full-rank matrices	24

4.1.2	70% of full-rank	24
4.1.3	30% of full-rank	24
4.1.4	10% of full-rank	24
4.1.5	5% of full-rank	25
4.1.6	Reduction in error	25
4.1.7	Screenshots	26
4.2	Paragraph vectors	26
4.2.1	Screenshots	26
4.2.2	gradient drop and loss drop	26
5	Conclusion	39
5.1	Data	39
6	Future Work	40
6.1	Correlation	40
6.2	Noise data removal	40
6.3	Get true value	40
	Bibliography	41

Part 1: Find similar Galaxy tools

1 Introduction

1.1 Galaxy

Galaxy¹ is an open-source biological data processing and research platform. It supports numerous types of extensively used biological data formats like FASTA, FASTAQ, GFF, PDB and many more. To process these datasets, Galaxy offers tools and workflows which either transform these datasets from one type to another or manipulate them. A simple example of data processing is to merge two compatible datasets to make one. Another example can be to reverse complement a sequence of nucleotides².

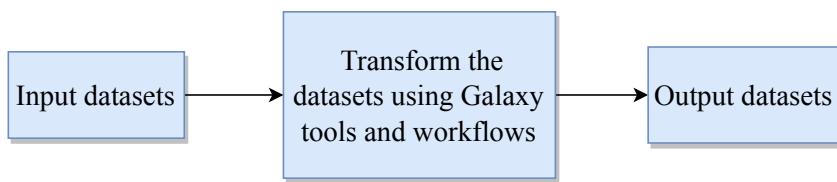


Figure 1: Basic flow of dataset transformation: it shows the basic flow of dataset transformation using Galaxy tools and workflows

A tool is a data-transforming entity which allows one or more types of datasets, transforms these datasets and produces output datasets. The tools are classified into multiple categories based on their functions. For example, the tools which manipulate text like replacing texts and selecting first lines of a dataset are grouped together under "Text Manipulation" group.

These tools form the building blocks of workflows. The workflows are data processing pipelines where a set of tools are joined one after another. The connected tools need to be compatible with each other which means the output types of one tool should be present in the input types of the next tool. A workflow can have one or more starting and end tools.

¹<https://usegalaxy.eu/>

²https://usegalaxy.eu/?tool_id=MAF_Reverse_Complement_1&version=1.0.1&__identifier=zmk9dx9ivbk

1.2 Galaxy tools

A tool entails a specific function. It consumes datasets, brings about some transformations and produces output datasets which can be fed to other tools. A tool has multiple attributes which include its input and output file types, name, description, help text and so on. They carry more information about a tool. When we look at the collective information about all these attributes for multiple tools, we find that some tools have similar functionalities based on their similarities in their corresponding attributes. For example, there are tools which share similarities in their respective functions and the input and output dataset types they are glued to. For example, a tool "hicexplorer hicpca"³ has an output type named "bigwig". Hence, if there is a tool or a set of tools which also has "bigwig" as their input and/or output type, we consider there could be some similarity between "hicexplorer hicpca" and the other tools as they do transformations on similar types of datasets. In addition, we can find similar functions of tools by analyzing their "name" and "description" attributes. Let's take an example of two tools (Figure 2):

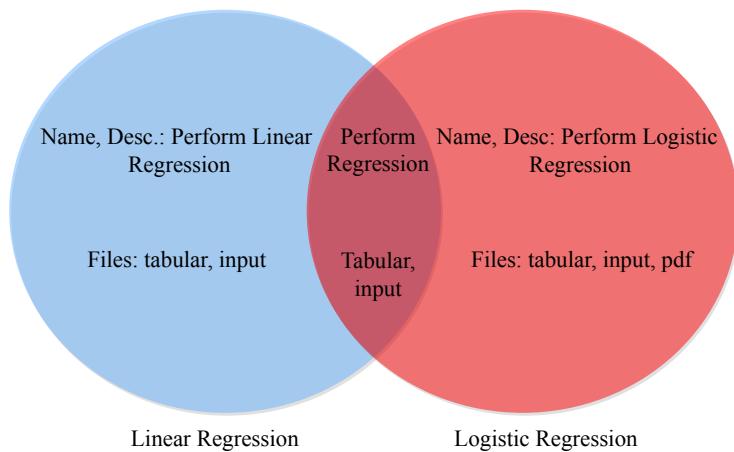


Figure 2: Venn diagram: it shows common features extracted from multiple attributes for the two tools

In figure 2, we take two tools - "Linear Regression" and "Logistic Regression" and collect their respective information from their input, output file types, name and description attributes. We see that these tools share features in the venn diagram. They both do regression and few file types are also common. In the same way, if we extrapolate this venn diagram and match one tool against all other tools, we hope

³https://usegalaxy.eu/?tool_id=toolshed.g2.bx.psu.edu/repos/bgruening/hicexplorer_hicpca/hicexplorer_hicpca/2.1.0&version=2.1.0&__identifier=5kcqmvb71gx

to find a set of tools similar in nature to the former tool. While searching for the related tools for a tool, it is possible that we end up with an empty set.

1.3 Motivation

From figure 2, we see that there can be tools which share characteristics. Galaxy has thousands of tools having a diverse set of functions. Moreover, new tools keep getting added to the older set of tools. From a user's perspective, it is hard to keep knowledge about so many tools. It is important to make a user aware of the presence of new tools added. If there is a model which dispenses a clue that there is a set of say n tools which are similar to a tool, it would give more options to a user for her/his data processing.

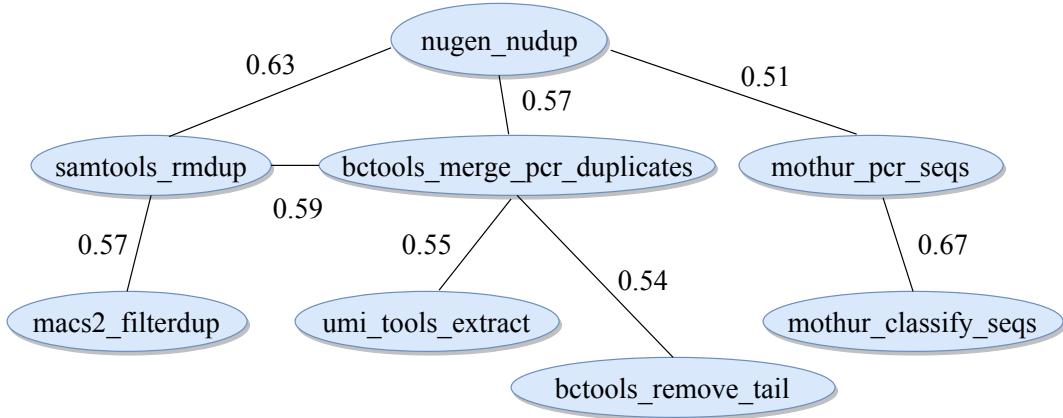


Figure 3: Similarity knowledge network: it shows how the tools in the network are related. The real numbers show the relation strength

To elaborate it more, let's take an example of a tool "nugen nudup"⁴. It is used to find and remove PCR duplicates. The similar tools for it can be "samtools rmdup" and "bctools merge pcr duplicates" which work on related concepts. These similar tools would further have their respective set of similar tools thereby making a network of related entities (tools). This "knowledge network" can help a user find multiple ways to process her/his data and exhibits "connectedness" among tools. The strength of this relation may vary from being small to large. To ascertain that, this study learns a continuous representation of the relation strength. Figure 2 shows how this knowledge graph can evolve. First, we find similar tools for "nugen nudup" and

⁴https://toolshed.g2.bx.psu.edu/repository?repository_id=4f614394b93677e3

connect them to their source tool specifying the similarity values as real numbers at the edges. These similar tools further have their own sets of similar tools and so on.

2 Approach

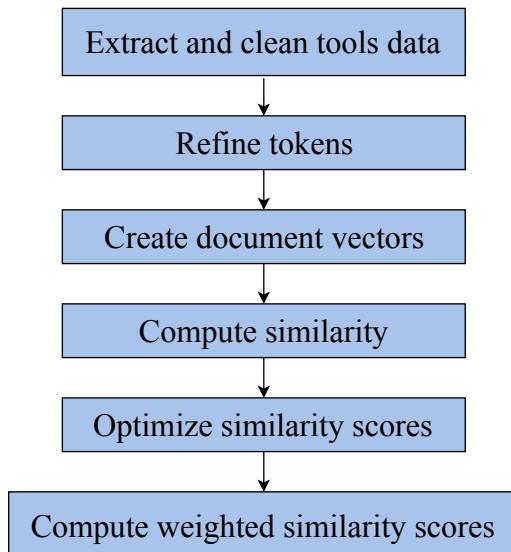


Figure 4: Flowchart to find similar tools: the flowchart highlights the important steps in finding similarity in tools.

2.1 Extract data

There are multiple repositories of Galaxy tools stored at GitHub¹. In each of the tool repository, there are *xml* files starting with a *< tool >* tag. We read all of these *xml* files, extract information from a few of the attributes and collect them in a *tabular* file.

2.1.1 Select tools attributes

A tool has multiple attributes like input and output file types, help text, name, description, citations and some more. But all of these attributes are not important and do not generally identify a tool exclusively. We consider some of these attributes:

¹One example:<https://github.com/galaxyproject/tools-iuc/tree/master/tools>

- Input and output file types
- Name and description
- Help text

Moreover, we combine the input and output file types and name and description respectively as they are of similar nature. These two combined attributes give complete information about a tool file types and its functionality. We also consider help text attribute which is larger in size compared to the previous two. At the same time, they are empty for few tools. Apart from being large in size, this attribute is noisy as well. It provides more information about the usage of a tool. Generally, in the first few lines, it gives a detailed explanation of the tool functions. Further, it explains how the input data should be supplied to a tool or how an input data looks like. Hence, much of the information present in this attribute is not important. Because of noise present in this attribute, we decide to use only upto first 4 lines which illustrates the core functionality of the tool. The decision to select only first 4 lines is empirical. The rest of the information in help text is discarded.

2.1.2 Clean data

Remove duplicates and stopwords

The collected data for tools is raw containing lots of commonplace and duplicate items which do not add value. These items should be removed to get *tokens* which are unique and useful. For example, a tool *bamleftalign* has input files as *bam*, *fasta* and output file as *bam*. While combining the file types, we discard the repeated file types and in this case, we consider file types as *bam*, *fasta*. The other attributes we deal with are different from the file types. The files types are discrete items but in attributes like name and description and help text, the account is in a human language. The explanation contains complete or partially complete sentences in *English*. Hence, to process this information, we need startegies that are prevalent in natural language processing². The sentences we write in *English* contain many words and has different parts. These parts include subject, object, preposition, interjection, verbs, adjectives, adverbs, articles and many others. For our processing, we need only those tokens (words) which categorize a tool uniquely and do away with multiple parts of speech present in the statements. For example, a tool named *tophat* has name and description as "TopHat for Illumina Find splice junctions using RNA-seq

²<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3168328/>

data". The words like *for*, *using* and *data* do not give much value as they must be present for many tools. These words are called as "stop words"³ and we selectively discard them. In addition, we remove numbers and convert all the tokens to lower case.

Use stemming

After removing duplicates and stop words, our data is clean and contain tokens which uniquely identify corresponding tools. When we frame sentences, we follow grammar which constrains us to use different forms of the same word in varying contexts. For example, a word *regress* can be used in multiple forms as *regresses* or *regression* or *regressed*. They share the same root and point towards the same concept. If many tools use this word in varying forms, it is beneficial to converge all the different forms of a word to one basic form. This is called stemming⁴. We use nltk⁵ package for stemming.

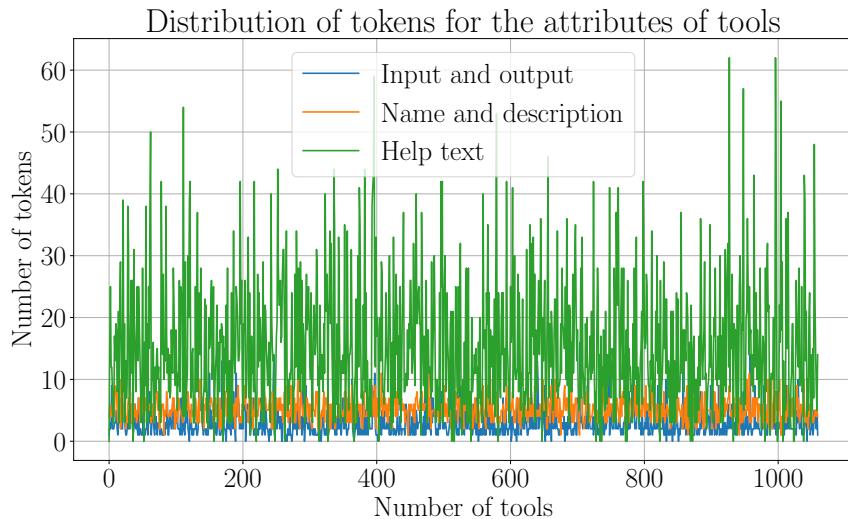


Figure 5: Distribution of tokens: the plot shows the distribution of tokens for all the attributes. The help text has the most number of tokens for all tools and input and output has the least number of tokens.

³<https://www.ranks.nl/stopwords>

⁴<https://nlp.stanford.edu/IR-book/html/htmledition/stemming-and-lemmatization-1.html>

⁵<http://www.nltk.org/>

Refine tokens

At this stage of tools preprocessing, we have a set of good tokens for each attribute which are input and output file types, name and description and help text. Let's call these sets as *documents*. The tokens present in these documents do not carry equal importance. Some tokens are more relevant to the document and some not so relevant. We need to find out importance factor for all tokens in a document. Using these factors, we can arrange them in big, sparse documents-tokens matrix. In this matrix, each row represents a document and each column belongs to one token. To compute these relevance scores, we use bestmatch²⁵. Let's associate some variables to be used in explaining this algorithm.

- Token frequency ⁶ tf
- Inverted document frequency idf
- Average document length $|D|_{avg}$
- Number of documents N
- Size of a document $|D|$

First of all, token frequency (tf) specifies the count of a token's occurrence in a document. If a token *regress* appears twice in a document, its tf is 2. This can also be understood as a weight given to this term. Inverted document frequency for a token is defined as:

$$idf = \log \frac{N}{df} \quad (1)$$

where df is the count of the documents in which this token is present and N is the total number of documents. If we randomly sample a document, then the probability of this token to be present in this document is $p_i = \frac{df}{N}$. From information theory, we can say that the information contained by this event is $-\log p_i$. The entity idf is higher when a token appears less number of documents which means that this token is a good candidate for representing the document and possesses higher power to distinguish between documents. The tokens which appear in many documents are not good representatives. Average document length is the average number tokens for all the documents. Size of a document is the count of all the tokens for that document [3].

⁶<https://nlp.stanford.edu/IR-book/pdf/06vect.pdf>

$$\alpha = (1 - b) + \frac{b \cdot |D|}{|D|_{avg}} \quad (2)$$

$$tf^* = tf \cdot \frac{k + 1}{k \cdot \alpha + tf} \quad (3)$$

$$BM25_{score} = tf^* \cdot idf \quad (4)$$

where k and b are hyperparameters. Using the equation 4, we compute the relevance score for each token in all the documents. Table 1 shows some sample scores for a few documents where the tokens are present with their respective relevance scores. In this way, we arrange document-tokens matrix for all the attributes of tools. For input and output file types, these matrix entries will have only two value, 1 if a token is present for a document and 0 if not. For other attributes, relevance scores are positive real numbers. This strategy of representing documents with their tokens is called vector space model as each document represents a vector of tokens.

Documents/tokens	regress	linear	gap	mapper	perform
LinearRegression	5.22	4.1	0.0	0.0	3.84
LogisticRegression	3.54	0.0	0.0	0.0	2.61
Tophat2	0.0	0.0	1.2	1.47	0.0
Hisat	0.0	0.0	0.0	0.0	0.0

Table 1: Document-tokens matrix: it stores relevance scores for each token for all the documents

Figure 4 shows the heatmaps for documents-tokens matrices that belong to input and output file types and name and description. We can see that these plots are sparse. Each entry in these matrices contain BM25 score for each token in every document. The representation shows how to find tokens which are good representatives of documents with a weighted by their relevance factors. But, they do not tell us anything about the co-occurrence of a few tokens in a document. It tells us that a token is important for a document if the BM25 score is higher but it does not tell us anything about its relation to other tokens. Due to this shortcoming, it does not acknowledge the presence of "concepts" or "context" hidden in a document. A concept in document can be realised when we see the relation among a few words. To illustrate this idea, let's take an example of three words - "New York City". These

three words mean little or point to different things if looked at separately. But, if we see them together, it points towards a concept. This vector space model lacks the ability to find the correlation among tokens. To enable the vector space model to learn this hidden concepts and find correlation among multiple tokens, we explore two ideas

- Latent Semantic Indexing/Analysis ⁷
- Paragraph Vectors

Using these approaches, we learn dense, n dimensional vector for each document instead of using sparse vectors as shown in figure 5.

Documents-tokens matrices for tools attributes

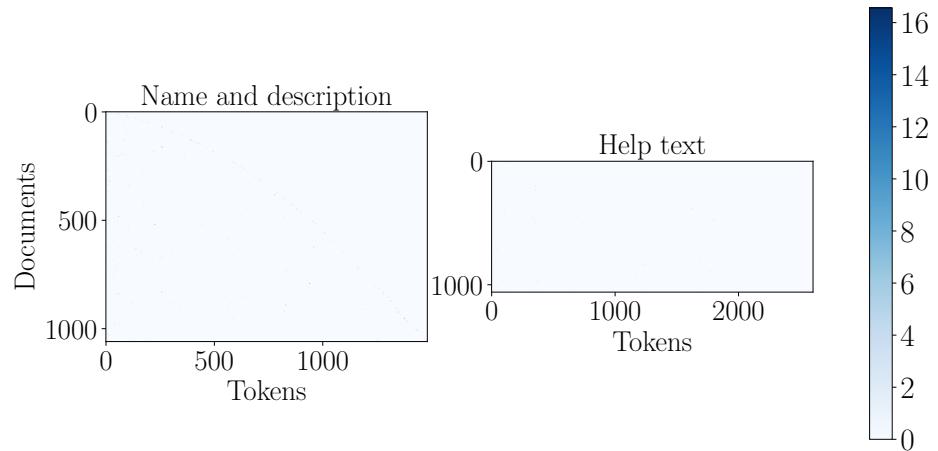


Figure 6: Documents-tokens matrices: the heatmap shows the documents-tokens sparse matrices for two attributes.

2.2 Document embeddings

2.2.1 Latent semantic indexing

It is statistical way to learn the hidden concepts in documents by computing a low-rank representation of a documents-tokens matrix. This low-rank matrix is dense

⁷<http://lsa.colorado.edu/papers/dp1.LSAintro.pdf>

(figure 6). We use singular value decomposition (*SVD*) to decompose the full-rank matrix into a significantly lower rank matrix. The optimal rank to which a matrix to be decomposed is empirical in nature. The rank chosen still contains most of the singular values. This decomposition follows the equation:

$$X_{n \times m} = U_{n \times n} \cdot S_{n \times m} \cdot V_{m \times m}^T \quad (5)$$

where n is the number of documents and m is the number of tokens. S is a diagonal matrix containing the singular values in descending order. It contains the weights of the concepts present in the matrix. The matrices U and V are orthogonal matrices which satisfy:

$$U^T \cdot U = I_{n \times n} \quad (6)$$

$$V^T \cdot V = I_{m \times m} \quad (7)$$

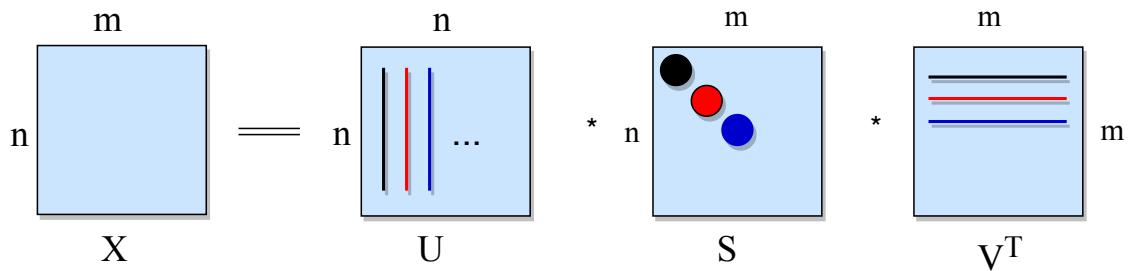


Figure 7: Pictorial representation of SVD: the figure shows how singular value decomposition is done.

Figure 6 explains ⁸ how *SVD* of a matrix is carried out. The matrix U contains information about how the tokens, arranged along the columns, are mapped to concepts and matrix V stores information about how the concepts are mapped to documents arranged along the rows.

Low-rank approximation

The low-rank approximation of a matrix is important to discard the features which are non-repeating or noise. Using this, we can collect the latent relations present in the documents-tokens matrices. We saw that our documents-tokens matrices suffer from sparseness and exhibit no relation among tokens. The approximation deals with

⁸<http://theory.stanford.edu/~tim/s15/l19.pdf>

these issues as well. The resulting matrices are dense and contain most of the singular values. The singular values which are small (the last entries of the s matrix along the diagonal) are discarded [4]. The low-rank approximated matrix X_k is computed as:

$$X_{n \times m} = U_k \cdot S_k \cdot V_k^T \quad (8)$$

where U_k is the first k columns of U , V_k is the first k rows and S is the first k singular values. k is an empirical parameter. X_k is called as the rank- k approximation of the full rank matrix X . Figure 7 shows the variation of the sum of singular values with the percentage rank of matrices. The percentage rank is $k \div K$ where $1 \leq k \leq K$ and K is the original (full) rank of a matrix. We take it as the original ranks of the three matrices are not comparable. By doing this, we can show them on one plot. Similar idea we use for the sum of the singular values of each matrix. From figure 7, we can say that if we reduce the ranks of matrices to 70% of the full-rank, we can still capture $\approx 90\%$ of the singular values. The reduction to half of the full-rank achieves $\approx 80\%$ of the sum of singular values.

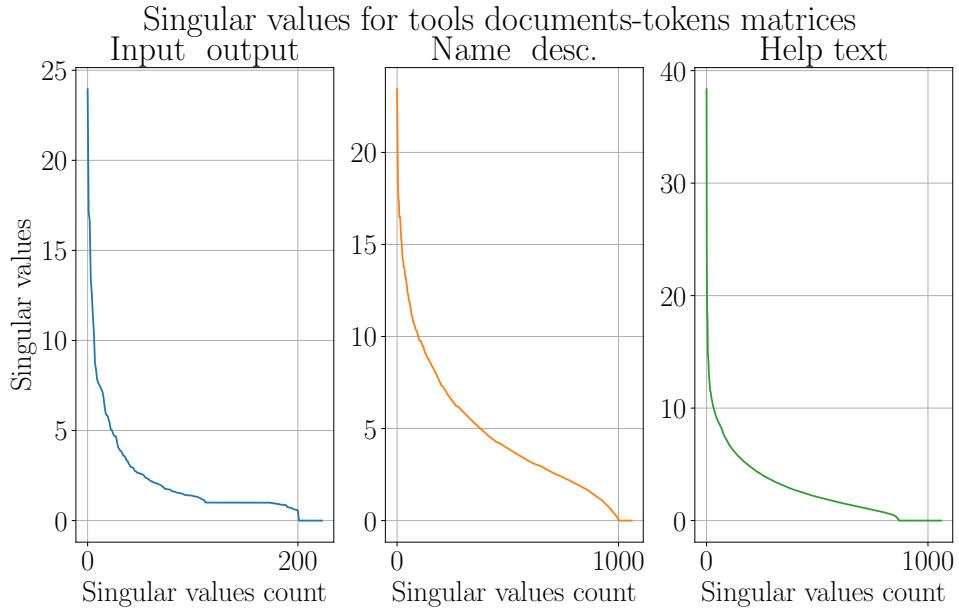


Figure 8: Matrix rank and singular values: the plot shows how the sum of singular values vary with the rank of the documents-tokens matrix for all the attributes.

We reduce the rank of the original documents-tokens matrices and compute the dense and approximated low-rank matrices. Figure 8 shows the low-rank matrices for

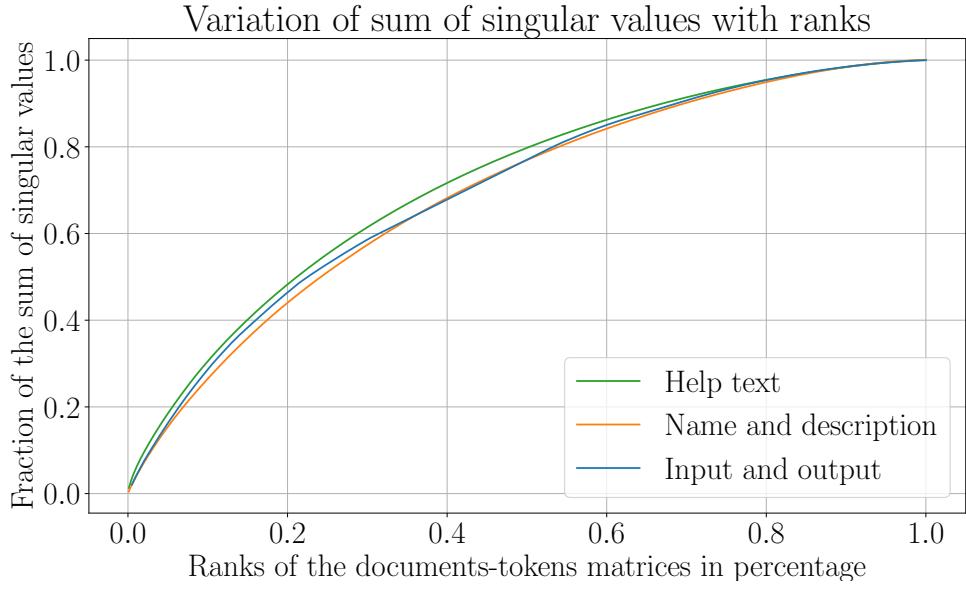


Figure 9: Matrix rank and singular values: the plot shows how the sum of singular values vary with the rank of the documents-tokens matrix for all the attributes.

input and output file types and name and description attributes. For computing this, we use only 40% of the full-rank. We can compare it with figure 5 and deduce that figure 8 is denser than figure 5. In these low-rank matrices, we have dense vector representations for documents along the rows. In each matrix, each row contains a vector for one document. Using these documents vectors, we can compute the correlation or similarity using any similarity measure. There are multiple similarity measures that can be used like euclidean distance, cosine similarity, manhattan distance. We use cosine angle similarity to compute the correlation between vectors. By this, we get a positive real number between 0.0 and 1.0 specifying how similar a pair of documents are. The higher the score, higher is the similarity. Computing this similarity for all the documents give us a similarity matrix $S_{n \times n}$ where n is the number of documents. This square, symmetric matrix is called as similarity or correlation matrix. We compute three such matrices each corresponding to one attribute.

Low rank representation of documents-tokens matrices

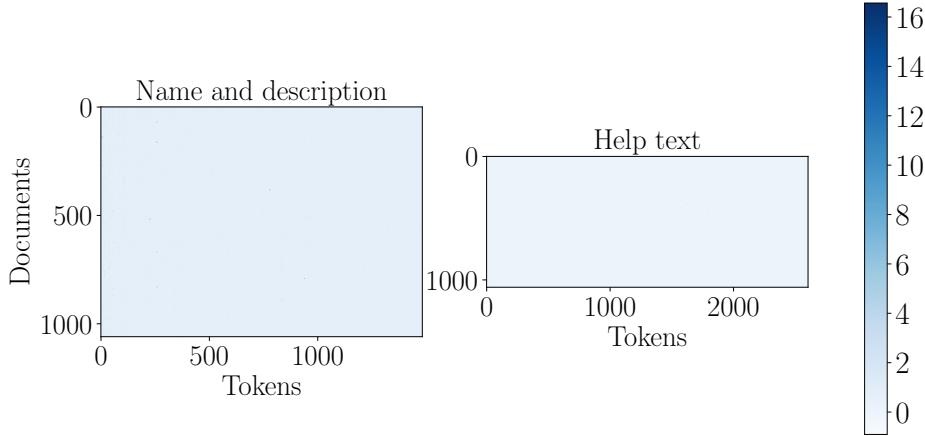


Figure 10: Documents-tokens matrices: the heatmap shows the documents-tokens low rank representation of the matrices.

2.2.2 Paragraph vectors

Using latent semantic indexing (LSI), we learnt dense vectors to represent each document. It learns better vector representations for documents compared to using full-rank documents-tokens matrices. One main limitation is to assess the quantity by which we need to lower the rank of a matrix in order to find the optimal results. There are ways to find the optimal reduced rank optimizing the frobenius norm but it is not simple. We would see in the analysis section that similar tools are more dominated by the scores shared by the input and output file types due to which the tools which similar in their functions do not come up. In order to avoid these limitations, we use an approach known as *doc2vec* (document to vector). It learns a dense, fixed-size vectors for each document using neural networks. These vectors are unique in a way that captures the semantics present in the documents. The documents which share similar context are represented by similar vectors. When we compute cosine distance between these vectors sharing similar context, we get a higher number (close to 1.0). It allows the documents to have variable lengths.

Approach

Paragraph vectors approach learns vector representations for all the words present in a set of documents. The words which are used in a similar context have similar vectors. For example, words like "data" and "dataset" which are used, in general, in similar context have are represented by close vectors. The vector representations of words in a corpus is learnt by finding the maximum of the following equation:

$$\frac{1}{T} \cdot \sum_{t=k}^{T-k} \log p(w_t | w_{t-k}, \dots, w_{t+k}) \quad (9)$$

where T is the total number of words in a corpus, k is the window size. We take a few words which make a context and using this context we try to predict each word. The probability p is computed using a softmax classifier and backpropagation is used to compute the gradient and the vectors are optimized using stochastic gradient descent. To learn paragraph vectors, in addition to using words vectors, paragraph vectors are also used to learn probability of the next words in a context. The paragraph and word vectors are averaged to make the classifier which predict next words in a context. There are two way how to choose the context.

- Distributed memory: In this approach of learning paragraph vectors, fixed length window of words are chosen and paragraph vectors and word vectors are used to predict the words in this context. The words vectors are shared across all paragraphs and paragraph vector is unique to each paragraph.
- Distributed bag of words: The words are randomly extracted from paragraphs and in this set of words, a word is chosen randomly and predicted using the paragraph vectors. The order of the randomly chosen words is ignored.

The figures 9 and 10 are inspired from the original work - Distributed Representations of Sentences and Documents⁹. The second form of learning paragraph vectors is simple and we use it to learn documents (paragraphs) vectors for name and description and help text attributes. We learn only the paragraph vectors which makes it less computationally expensive [5].

⁹<https://arxiv.org/abs/1405.4053>

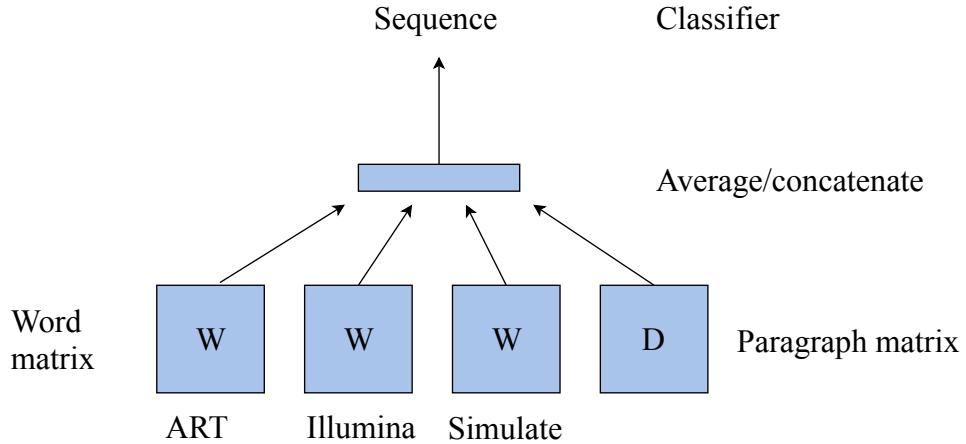


Figure 11: Paragraph vectors distributed models: it shows the mechanism how the paragraph and words vectors contributed to predict next words in a context.

2.3 Similarity measures

We have vectors representing the documents for all the three attributes we consider to compute similar tools. To find the similarity in a pair of vectors, we need to apply a similarity measure to get a similarity score. This score quantifies how much similar a pair of documents are which means how close or distant these documents are.

2.3.1 Cosine similarity

It gives the cosine angle value between a pair of documents vectors. Let's say we have two vectors, x and y . We can write:

$$x \cdot y = |x| \cdot |y| \cdot \cos \theta \quad (10)$$

$$\cos \theta = \frac{x \cdot y}{|x| \cdot |y|} \quad (11)$$

where $|.|$ is the norm of the vector x . If the norm is 0, we use 0 for the value of $\cos \theta$. The values emitted by this similarity follows a natural progression that if the documents are dissimilar, then it is 0 and if completely similar, it is 1.0. Otherwise it lies between 0.0 and 1.0. This score can also be understood as a kind of probability of similarity between a pair of documents ¹⁰.

¹⁰<https://nlp.stanford.edu/IR-book/html/htmledition/dot-products-1.html>

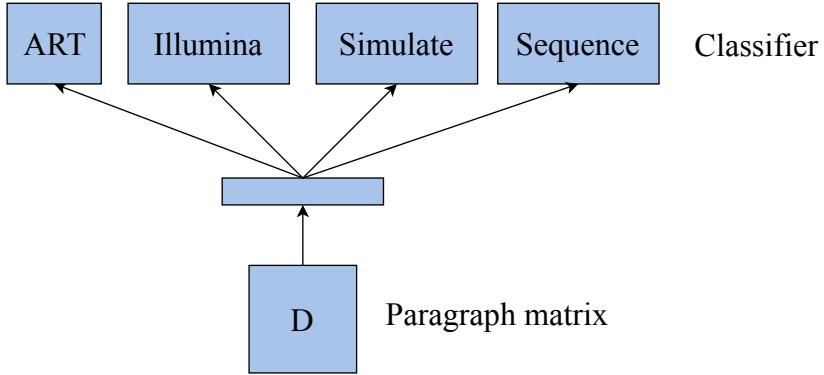


Figure 12: Distributed bag of words: it shows how a paragraph vector is used to predict a randomly chosen word.

2.3.2 Jaccard index

2.4 Optimization

Using the paragraph vectors and applying cosine similarity, we get similarity or correlation matrices one each for input and output file types, name and description and help text attributes. The matrices have same dimensions ($n \times n$, n is the number of tools). To combine these matrices, one simple idea would be to take an average of the corresponding rows of scores from the matrices for a document. This would generate a matrix of the same dimension where each row would correspond to one document. This row would contain similarity scores of a document against all other documents. The diagonal entries of this matrix would be 1.0. All other entries would be a positive real number between 0.0 and 1.0. Another way to find the combination is to learn the weights on the rows from three matrices and them combine them to obtain optimal similarity scores for a tool. The weight would be a positive real number between 0.0 and 1.0. Instead of using a fixed importance factor (weights) of $1/3$ (3 is the number of matrices), we use optimization to find these real numbers and then combine the matrices by multiplying with these weights.

$$S_k^{optimal} = w_{io}^k \cdot S_{io}^k + w_{nd}^k \cdot S_{nd}^k + w_{ht}^k \cdot S_{ht}^k \quad (12)$$

where w is a positive, real number and satisfy $w_{io}^k + w_{nd}^k + w_{ht}^k = 1$. S_{io}^k , S_{nd}^k and S_{ht}^k are the similarity scores (corresponding matrix rows) for k^{th} tool corresponding to input and output file types, name and description and help text attributes respectively.

We have these similarity scores. But we need to learn these importance weights.

We use gradient descent optimizer to learn these weights against a loss function. If we look closely the similarity measure, we see that the maximum similarity between a pair of documents can be 1.0. Hence, the ideal similarity scores for one document against all other documents:

$$S_{ideal} = [1.0, 1.0, \dots, 1.0] \quad (13)$$

where S_{ideal} is the ideal similarity scores for one document against all other document. Using this ideal scores, we can compute the error we accrue for all the similarity scores from three attributes. After computing the error, we can verify which attribute is closer to the ideal score and which are not. The attributes which exhibit lower error get higher weights and those which score higher error get lower weights. The next section elaborates the online way to do the optimization.

Gradient descent

Gradient descent is a popular algorithm for optimizing an objective function with respect to its parameters. The parameters are the entities which we want to learn. In our case, these are the weights. The algorithm minimizes an error function. Let's frame the error function:

$$Error_{io}^k(w_{io}) = \sum_{t=1}^{N-1} \cdot (w_{io} \cdot S_{io}^t - S_{ideal})^2 \quad (14)$$

$$Error_{nd}^k(w_{nd}) = \sum_{t=1}^{N-1} \cdot (w_{nd} \cdot S_{nd}^t - S_{ideal})^2 \quad (15)$$

$$Error_{ht}^k(w_{ht}) = \sum_{t=1}^{N-1} \cdot (w_{ht} \cdot S_{ht}^t - S_{ideal})^2 \quad (16)$$

$$Error^k(w) = \sum_{t=1}^{N-1} \cdot (w \cdot S^t - S_{ideal})^2 \quad (17)$$

$$\arg \min_w Error^k(w) \quad (18)$$

where $Error$ is a vector of $\langle Error_{io}, Error_{nd}, Error_{ht} \rangle$, w is $\langle w_{io}, w_{nd}, w_{ht} \rangle$ and S is $\langle S_{io}, S_{nd}, S_{ht} \rangle$.

To minimize the equation 17, we need to compute the gradient. The gradient specifies the rate of change of error with respect to the weights.

$$\text{Gradient}^k(w) = \frac{\partial \text{Error}^k}{\partial w} = 2 \cdot \sum_{t=1}^{N-1} (w \cdot S^t - S_{ideal}) \cdot S^t \quad (19)$$

Using the gradient vector, we update the weights. For this we need to set the learning rate.

$$w^k = w^k - \eta \cdot \text{Gradient}^k \quad (20)$$

To find the right learning rate is important as a high value can diverge the gradient learning and a low value can slow down the learning. For each iteration of gradient descent, we employ backtracking line search to find the right learning rate which assists in finding the update of the weights parameters.

Backtracking line search

We use backtracking line search method to find the necessary decrease along the gradient search direction. It enables us to avoid all the small steps generally taken by gradient descent to decrease the objective error function. At each step of gradient descent, we start with 1.0 as learning rate and gradually decrease it until the following condition is met or gradient decrease stabilizes with a tolerance factor ϵ . The decrease of learning rate is controlled by a hyperparameter β . The name backtracking comes from the fact that it starts with unit learning rate and gradually decreases the learning rate.

$$\text{Error}(w^*) \approx \text{Error}(w) - \eta \cdot |\text{Gradient}(w)|^2 < \text{Error}(w) - \alpha \cdot \eta \cdot |\text{Gradient}(w)|^2 \quad (21)$$

where $w^* = w - \eta \cdot \text{Gradient}(w)$ and α lies in the range $(0, 0.5)$ ¹¹.

¹¹http://users.ece.utexas.edu/~cmcaram/EE381V_2012F/Lecture_4_Scribe_Notes.final.pdf

3 Experiments

3.1 Amount of help text

We use a maximum of 4 lines of text from the help text attribute while reading the xml files of tools. This attribute is noisy and contains text which is not useful to be set up as a basis for finding similarity. We need to be careful of the amount of text which we extract from this attribute. Experimentally, we verify that using 4 lines of text from this source works better in most of the cases.

3.2 One and two sources instead of three

We have used three different attributes for compiling the collection of tokens. These attributes are different from each other. Using them together to make one set of tokens for each tool would not be beneficial. We compute similarities for these different sources and combine them using optimization by learning optimal weights on each attribute.

3.3 Similarity measures

For finding similarity using input and output file types, we use jaccard index because we do not need to learn any "concept" hidden in a set of file types for a tool. For similarities computed for name and description and help text, we use cosine similarity. Both these similarity measures give a real number between 0.0 and 1.0 as a similarity score.

3.4 Latent Semantic Analysis

Using latent semantic analysis, we learn dense vector representation for a document. It reduces the rank of the document-tokens matrix. We need to find out this reduction factor which can improve the similarity scores among tools compared to using full-rank

document-tokens matrix. We follow an approach of lowering the rank by certain factors and look at how the values in these matrices are spread. We reduce the rank to 70%, 50%, 30%, 10% and 5% of the full-rank value. Moreover, we also verify the similarity matrices corresponding to these low-rank matrices representations. We expect the documents-tokens and similarity matrices to be more dense compared to no rank reduction. To reduce the ranks, we consider documents-tokens matrices of name and description and help text and leave the input and output matrix in its full-rank state. We singular value decomposition routine from *numpy* linear algebra package.

3.5 Paragraph Vector

In this approach, we learn a fixed-length dense vector representations for each document. We set the dimensions of these vectors. When the size of a document is lower, we use a lower number for computing the fixed-length vector. For example, in figure 5 we see that the number of tokens is higher for help text compared to name and description. To learn fixed-length vectors, we set the vector's length to be 50 for name and description and 300 for help text. Here as well, we learn paragraph vectors only for name and description and help text and not for input and output file types. We use *gensim* model to learn these paragraph vectors.

3.5.1 Distributed bag-of-words

We use this approach to learn vectors for each document. It does not use word-vectors and rely on paragraph vectors to predict the randomly chosen words from a sampled text window. Learning only the paragraph vectors is faster and computationally less expensive.

3.6 Gradient Descent

To optimize the similarity scores from coming from multiple attributes, we use gradient descent optimizer to learn weights on the scores of these attributes. Based on the similarity measures, we construct a error function to the optimizer which minimizes it.

3.6.1 Learning rates

We use backtracking line search to set learning rate for each iteration. It expedites the learning by computing a learning rate that sufficiently diminishes the error function in each iteration of gradient descent. We set a maximum iteration to 10 within which the learning saturates.

4 Results and Analysis

4.1 Latent Semantic Analysis

We experiment with multiple values of matrix rank reduction and we find that as we reduce the rank of documents-tokens matrices, they become more dense. As we reduce the ranks, the distribution of the learnt weights also change.

4.1.1 Full-rank matrices

Figure 13 shows the similarity matrices computed using full-rank documents-tokens matrices for name and description and help text. We do not apply rank reduction on documents-tokens matrix of input and output file types. Using these similarity matrices, we find an optimal combination by optimizing them. The "optimal" correlation plot shows this weighted average. Figure 14 shows the distribution of weights for multiple tool attributes. We see the magnitude of weights learnt for input and output file types is higher than the other two. It is associated with the higher values captured for the similarity matrix of input and output file types.

4.1.2 70% of full-rank

We reduce the ranks of two documents-tokens matrices to 70% of the original ranks. For example, if the rank of a matrix is 100, we reduce the rank to 70 using singular value decomposition. From figure 15, we see that the name and description and help text similarity matrices start becoming more dense compared to figure 13. The distribution of the weights also change (figure 16).

4.1.3 30% of full-rank

We reduce the ranks of two documents-tokens matrices to 30% of the original ranks.

4.1.4 10% of full-rank

We reduce the ranks of two documents-tokens matrices to 10% of the original ranks.

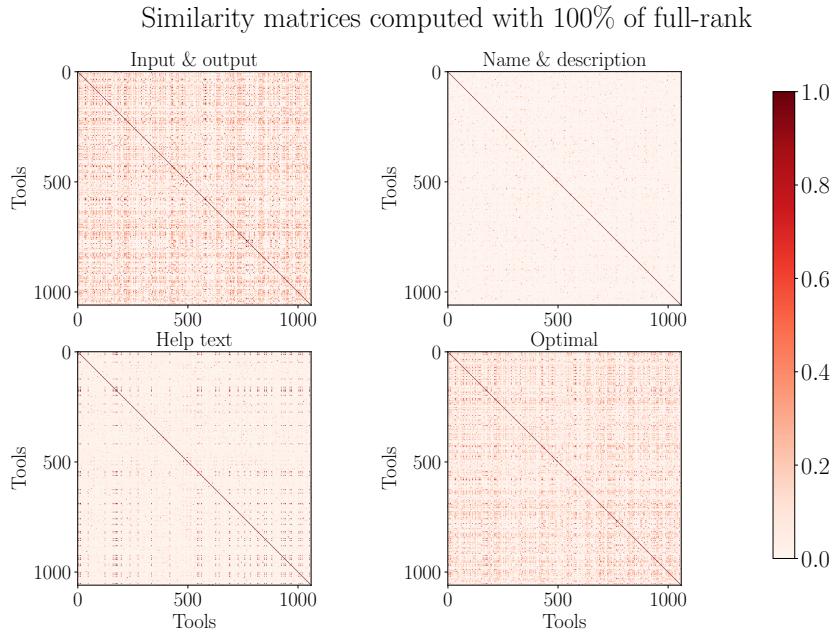


Figure 13: Similarity matrices using full rank: this correlation plot shows the individual similarity matrices for multiple attributes and the weighted average of these matrices as optimal one.

4.1.5 5% of full-rank

We reduce the ranks of two documents-tokens matrices to 5% of the original ranks.

4.1.6 Reduction in error

We observe the loss in mean squared error when we decrease the ranks of the matrices. When we reduced the ranks, the similarity scores for the name and description and help text increase. This increase accounts for their higher importance weights learnt by the optimizer. The importance weights on an average become more balanced and can account for the decrease in the mean squared error. We would also see that in the next section that the ranking of similar tools also become better after rank reduction.

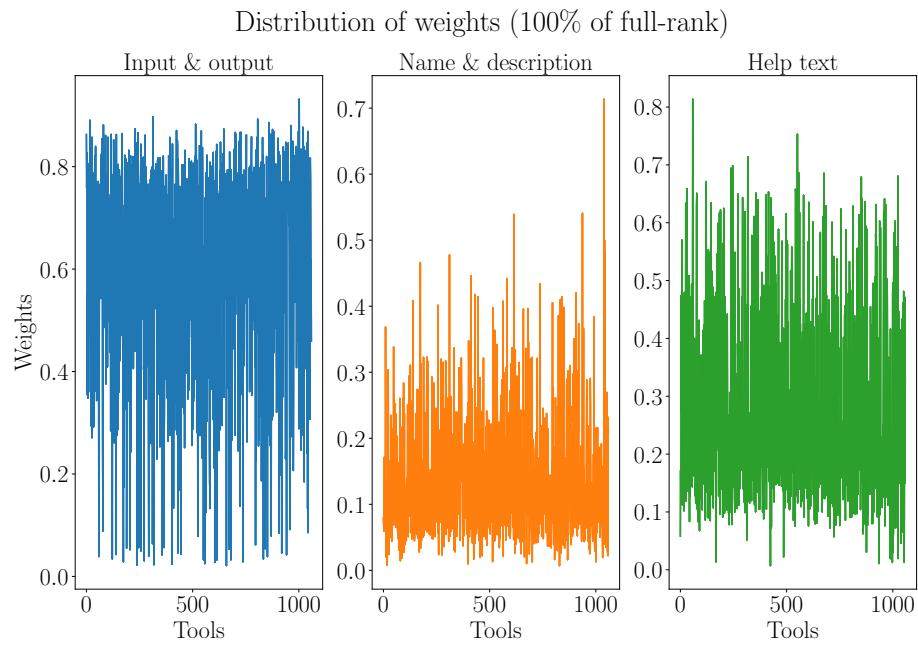


Figure 14: Weights distribution full rank: this plot show how the weights vary for multiple tool attributes.

4.1.7 Screenshots

4.2 Paragraph vectors

4.2.1 Screenshots

4.2.2 gradient drop and loss drop

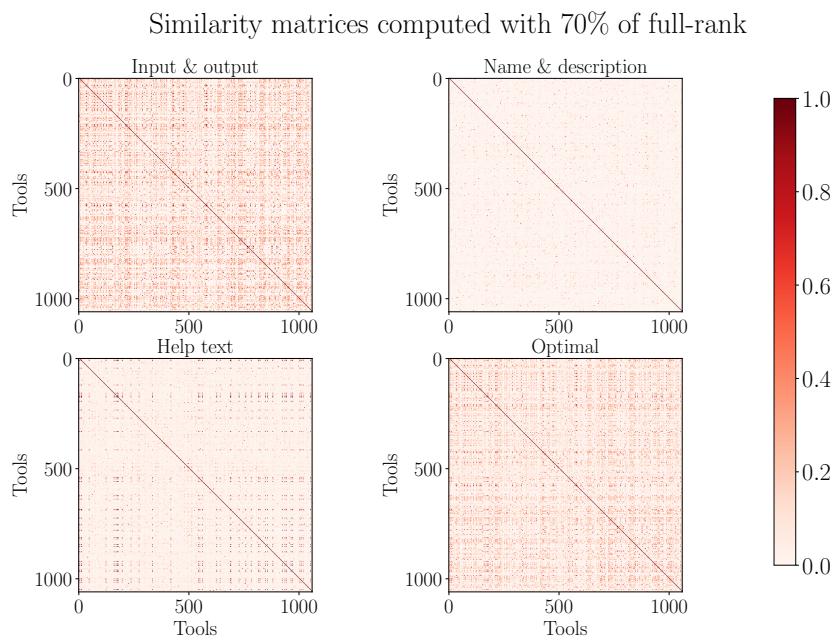


Figure 15: Similarity matrices using 70% rank: this correlation plot shows the individual similarity matrices for multiple attributes and the weighted average of these matrices as optimal one after reducing the name and description and help text matrices to 70% of their respective ranks.

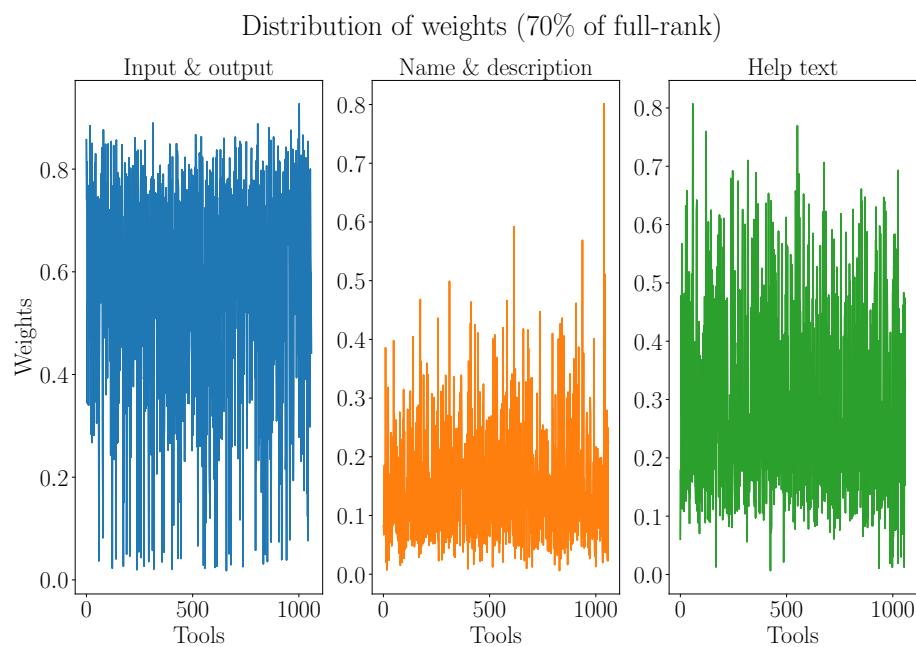


Figure 16: Weights distribution using 70% rank: this plot show how the weights vary for multiple tool attributes after reducing the name and description and help text matrices to 70% of their respective ranks.

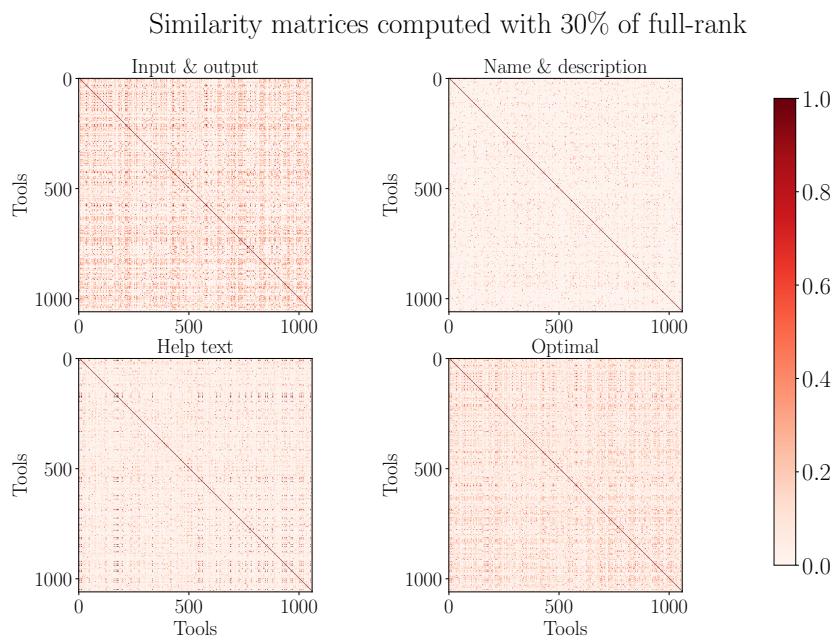


Figure 17: Similarity matrices using 30% rank: this correlation plot shows the individual similarity matrices for multiple attributes and the weighted average of these matrices as optimal one after reducing the name and description and help text matrices to 30% of their respective ranks.

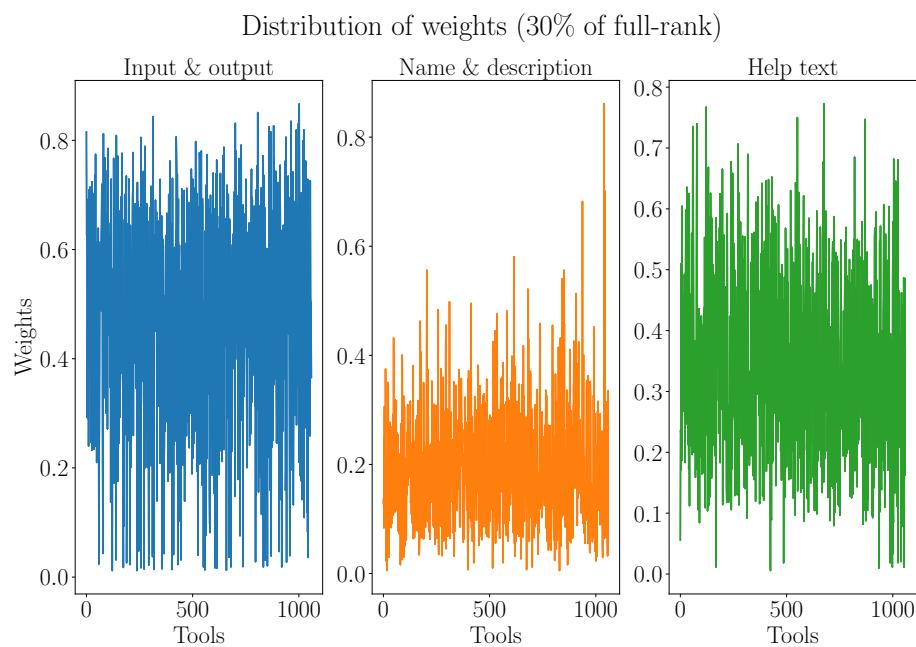


Figure 18: Weights distribution using 30% rank: this plot show how the weights vary for multiple tool attributes after reducing the name and description and help text matrices to 30% of their respective ranks.

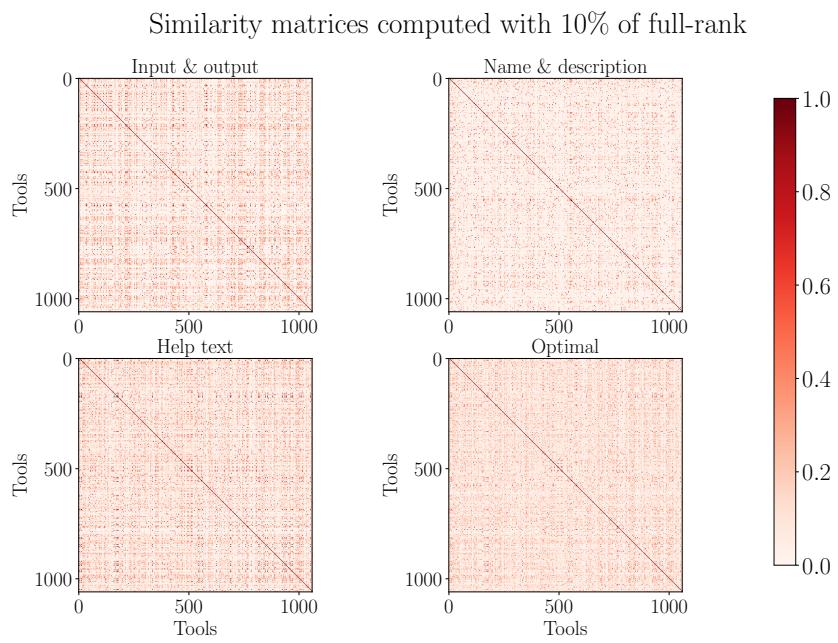


Figure 19: Similarity matrices using 10% rank: this correlation plot shows the individual similarity matrices for multiple attributes and the weighted average of these matrices as optimal one after reducing the name and description and help text matrices to 10% of their respective ranks.

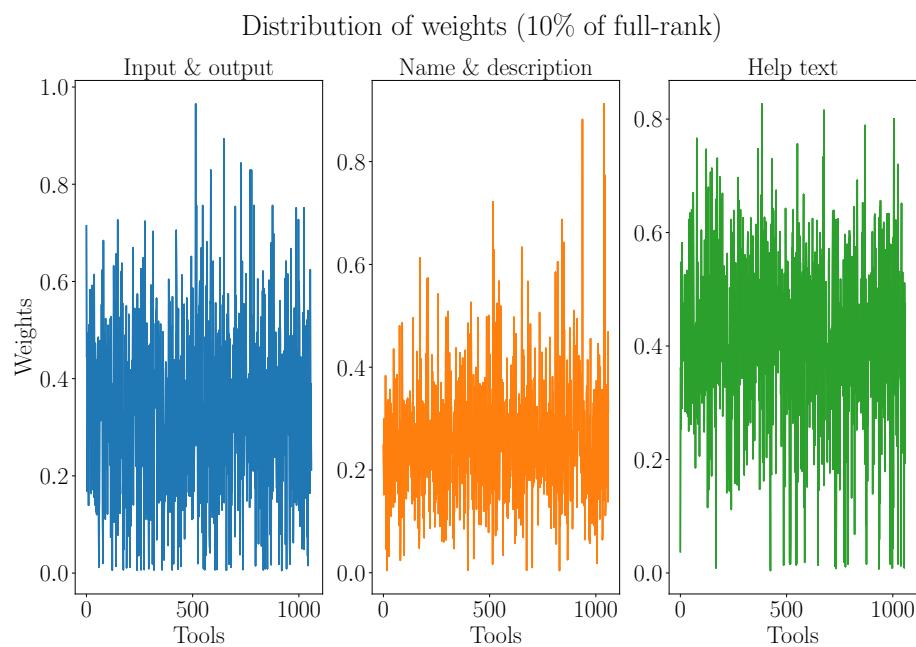


Figure 20: Weights distribution using 10% rank: this plot show how the weights vary for multiple tool attributes after reducing the name and description and help text matrices to 10% of their respective ranks.

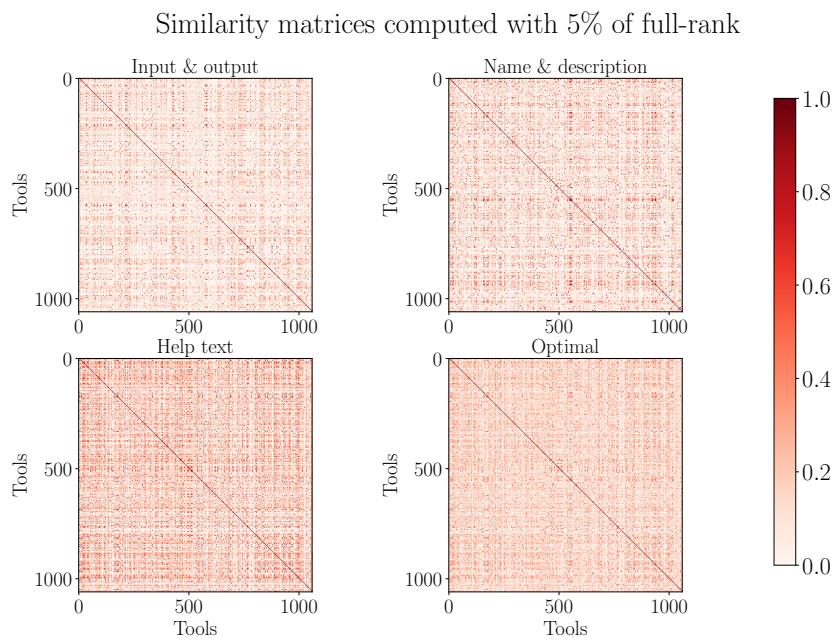


Figure 21: Similarity matrices using 5% rank: this correlation plot shows the individual similarity matrices for multiple attributes and the weighted average of these matrices as optimal one after reducing the name and description and help text matrices to 5% of their respective ranks.

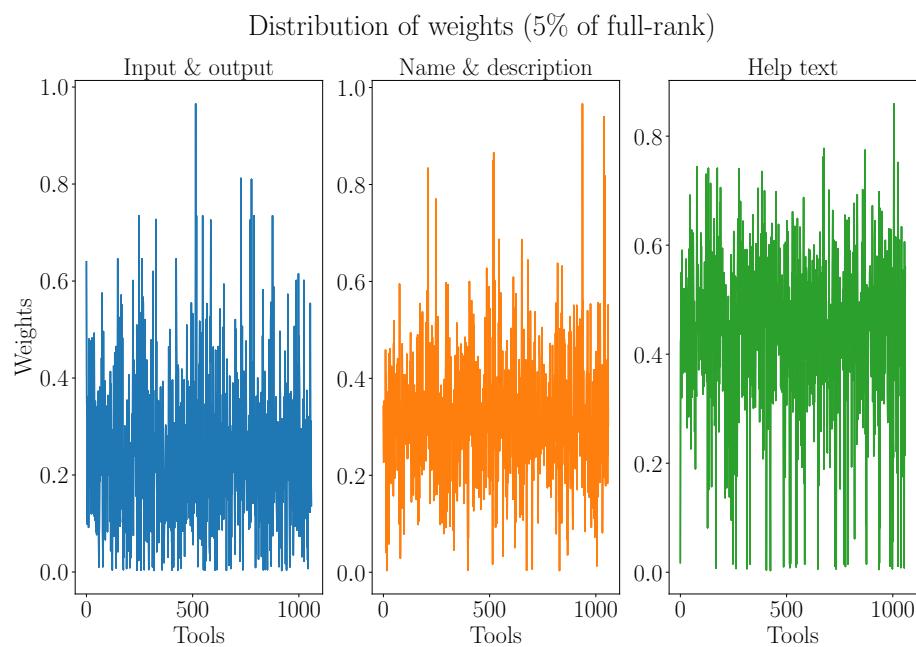


Figure 22: Weights distribution using 5% rank: this plot show how the weights vary for multiple tool attributes after reducing the name and description and help text matrices to 5% of their respective ranks.

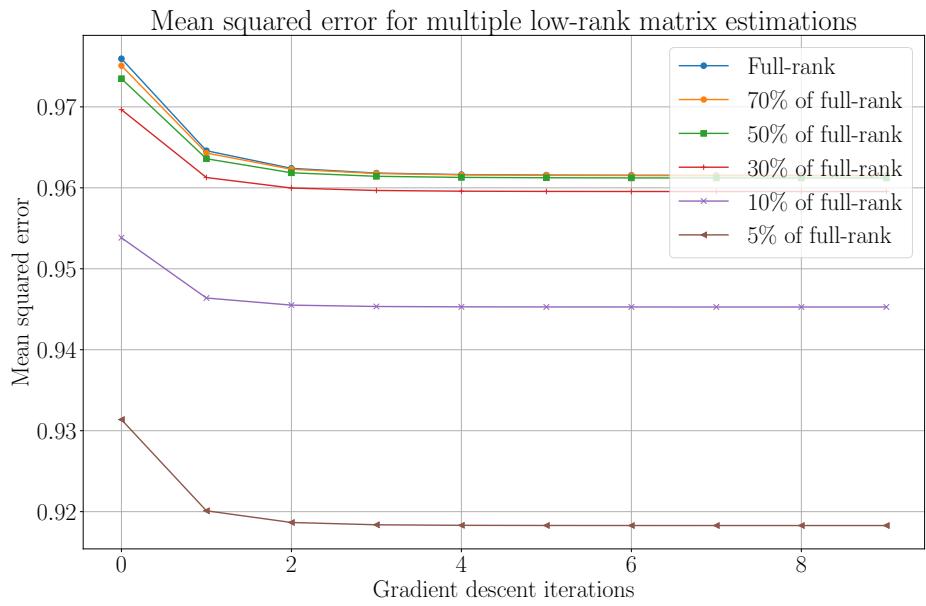


Figure 23: Similarity matrices using 5% rank: this correlation plot shows the individual similarity matrices for multiple attributes and the weighted average of these matrices as optimal one after reducing the name and description and help text matrices to 5% of their respective ranks.

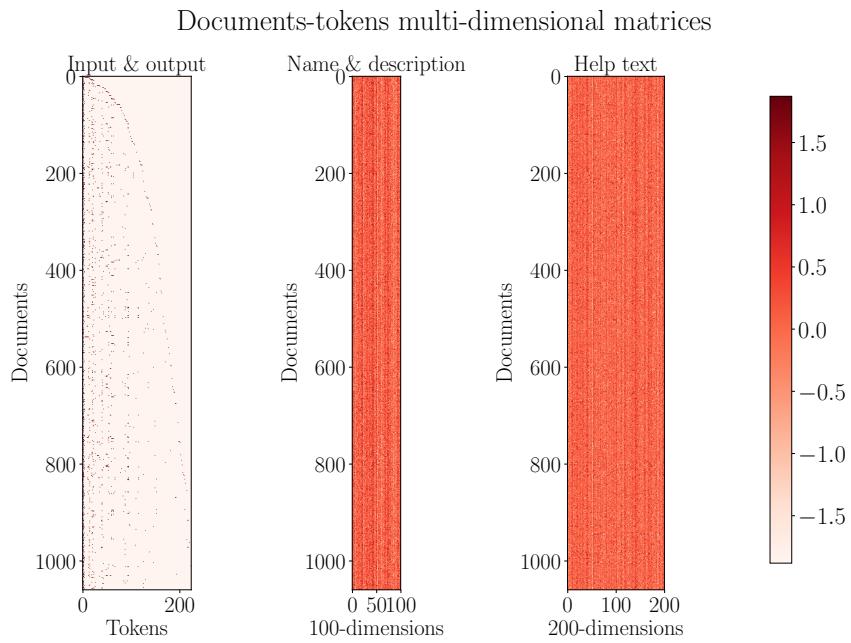


Figure 24: Documents-tokens matrices for Paragraph vectors approach: this heatmap shows a documents-tokens matrix (the leftmost) for input and output file types. The next two matrices belong to name and description and help text attributes. Each row of both of these matrices is a n-multi-dimensional paragraph vector.

Similarity matrices for multiple tools attributes and an optimal one

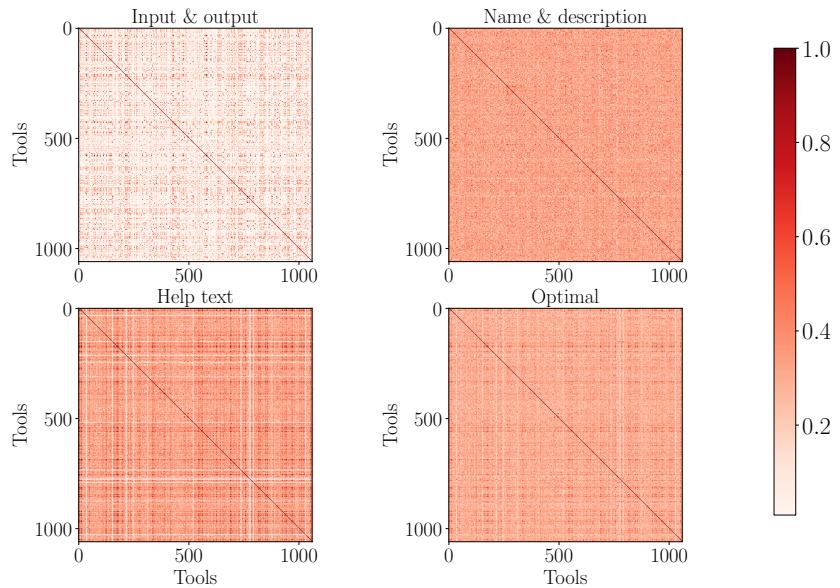


Figure 25: Similarity matrices computed using "paragraph vectors" approach: this plot show four correlation heatmaps. The top-left, top-right and bottom-left are the individual correlation matrices for input and output, name and description and help text attributes. The bottom-right one is the weighted average correlation matrix computed from the last three matrices.

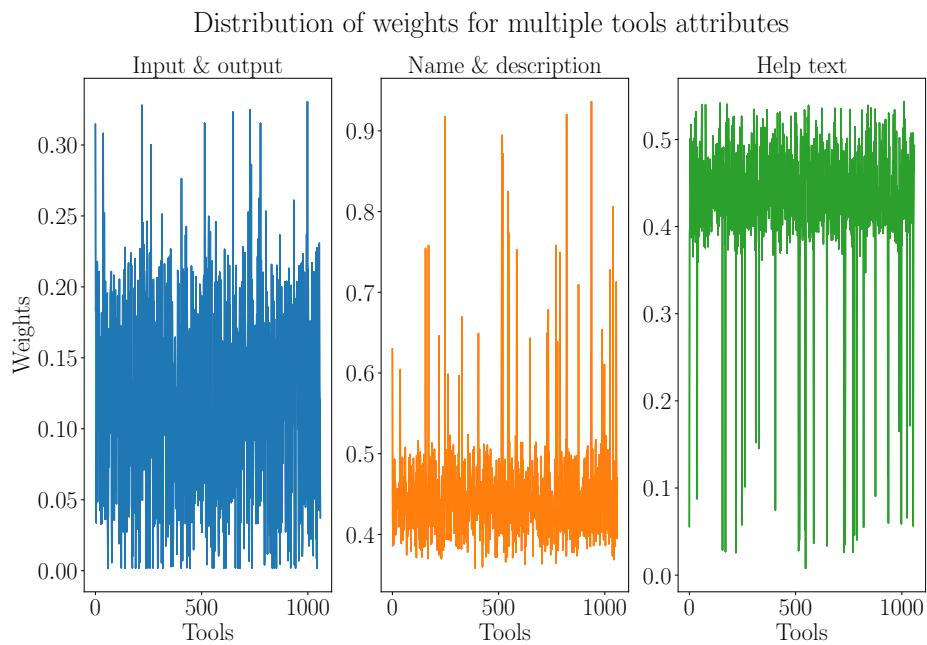


Figure 26: Weight distribution learnt using "paragraph vectors" approach: this plot shows the distribution of weights for multiple tools attributes learnt using "paragraph vectors" approach. We can see that the magnitude of weights for input and output file types are lower compared to two other attributes.

5 Conclusion

5.1 Data

Noisy, no true value, some good, some bad results.

6 Future Work

6.1 Correlation

6.2 Noise data removal

6.3 Get true value

Part 2: Predict next tools in Galaxy workflows

Bibliography

- [1] Z. C. Lipton, D. C. Kale, C. Elkan, and R. C. Wetzel, “Learning to diagnose with LSTM recurrent neural networks,” *CoRR*, vol. abs/1511.03677, 2015.
- [2] H. Sak, A. W. Senior, and F. Beaufays, “Long short-term memory based recurrent neural network architectures for large vocabulary speech recognition,” *CoRR*, vol. abs/1402.1128, 2014.
- [3] S. Robertson and H. Zaragoza, “The probabilistic relevance framework: Bm25 and beyond,” *Found. Trends Inf. Retr.*, vol. 3, pp. 333–389, Apr. 2009.
- [4] J. Yang, “Notes on low-rank matrix factorization,” *CoRR*, vol. abs/1507.00333, 2015.
- [5] Q. V. Le and T. Mikolov, “Distributed representations of sentences and documents,” *CoRR*, vol. abs/1405.4053, 2014.

